

Predicting Divorce with the DS Test: a quantitative approach

CYO - Data Science: Capstone

Luz E Rodriguez - October 2020

```
if(!require(magrittr)) install.packages("magrittr",
                                         repos = "http://cran.us.r-project.org")
if(!require(utils)) install.packages("utils",
                                       repos = "http://cran.us.r-project.org")
if(!require(ggplot2)) install.packages("ggplot2",
                                       repos = "http://cran.us.r-project.org")
if(!require(dplyr)) install.packages("dplyr", repos = "http://cran.us.r-project.org")
if(!require(kableExtra)) install.packages("kableExtra",
                                          repos = "http://cran.us.r-project.org")
if(!require(knitr)) install.packages("knitr",
                                      repos = "http://cran.us.r-project.org")
if(!require(tibble)) install.packages("tibble",
                                       repos = "http://cran.us.r-project.org")
if(!require(corrplot)) install.packages("corrplot",
                                         repos = "http://cran.us.r-project.org")
if(!require(caret)) install.packages("caret",
                                       repos = "http://cran.us.r-project.org")
if(!require(mlbench)) install.packages("mlbench",
                                       repos = "http://cran.us.r-project.org")
if(!require(gridExtra)) install.packages("gridExtra",
                                         repos = "http://cran.us.r-project.org")
if(!require(ggthemes)) install.packages("ggthemes",
                                         repos = "http://cran.us.r-project.org")
if(!require(factoextra)) install.packages("factoextra",
                                          repos = "http://cran.us.r-project.org")
if(!require(rpart)) install.packages("rpart",
                                      repos = "http://cran.us.r-project.org")
if(!require(rpart.plot)) install.packages("rpart.plot",
                                          repos = "http://cran.us.r-project.org")
if(!require(randomForest)) install.packages("randomForest",
                                             repos = "http://cran.us.r-project.org")
if(!require(Rtools)) install.packages("Rtools",
                                       repos = "https://cran.rstudio.com/bin/windows/Rtools/")
```

1. Project Overview

I found this interesting article about Divorce Prediction <https://dergipark.org.tr/en/download/article-file/748448>. This is based on a Gottman's couples therapy research of more than 20 years and proposes that couples can find alert signs of divorce. If diagnosed on time, marriages can work on tackling their problems to have a happy marriage and avoid divorce. Gottman has designed therapies to help marriages to improve their quality of life and relationships.

The article analyzes a sample survey in Turkey and predicts divorce with 98% accuracy by using correlation based feature selection and artificial neural networks.

It is very exciting to be able to quantify emotions and behaviors and apply analytics in social sciences. It is also more difficult and biased because there are a lot of assumptions, like people are completely honest and that every person interpret the questions in the same way. That is why I want to explore this dataset and apply other prediction methods to see if I get the same prediction rate. I will predict a categorical variable, in this case whether a person will remain married(0) or will get divorced(1) based on their current behavior explained by the answers to the survey.

2. Analysis

Data Set Description

The dataset is taken from UCI Machine Learning Repository - Center for Machine Learning And Intelligent Systems <https://archive.ics.uci.edu/ml/machine-learning-databases/00497/>. This corresponds to a dataset from a questionnaire that ask participants several questions related to their behavior with their partners based on the Divorce Predictors Scale (DPS) from Gottman couples therapy.

The initial research <https://dergipark.org.tr/en/download/article-file/748448> was developed by Dr. Mustafa Kemal Yöntem, Dr. Kemal ADEM, Prof. Dr. Tahsin İlhan, and Lecturer Serhat Kılıçarslan, inspired by Gottman's DPS couples therapy.

The survey had a total of 170 Turkish participants, 84 males and 86 females between 20 to 63 years old, from which 84 are divorced(49%) and 86 are married(51%) who answered 54 questions. All responses were collected on a 5-point scale: (0=Never, 1=Seldom, 2=Averagely, 3=Frequently, 4=Always).

The last variable **Class** indicates marital status as 0 for married and 1 for divorced.

Downloading the dataset

```
current_dir <- getwd()
divorce <- utils::read.table(file.path(current_dir,"divorce.csv"),
                             sep = ";", header = TRUE)
```

```
print("Split between married/divorced in dateset")
```

```
## [1] "Split between married/divorced in dateset"
```

```
(table(divorce$Class))
```

```
##
##  0  1
## 86 84
```

```
print("Summary scale responses")
```

```
## [1] "Summary scale responses"
```

```
with(divorce, summary(Atr1))
```

```
##      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.
##    0.000   0.000   2.000   1.776   3.000   4.000
```

```
print("Dataset glimpse")
```

```
## [1] "Dataset glimpse"
```

```
tibble::glimpse(divorce)
```

```
## Rows: 170
## Columns: 55
## $ Atr1 <int> 2, 4, 2, 3, 2, 0, 3, 2, 2, 1, 4, 4, 3, 3, 3, 4, 4, 4, 3, 4, 4...
## $ Atr2 <int> 2, 4, 2, 2, 2, 0, 3, 1, 2, 1, 4, 4, 4, 4, 4, 4, 4, 4, 3, 4, 3...
## $ Atr3 <int> 4, 4, 2, 3, 1, 1, 3, 2, 1, 1, 4, 4, 3, 3, 3, 3, 3, 4, 4, 4, 3...
## $ Atr4 <int> 1, 4, 2, 2, 1, 0, 2, 2, 0, 1, 3, 3, 4, 4, 4, 2, 2, 3, 4, 3, 3...
## $ Atr5 <int> 0, 4, 1, 3, 1, 0, 1, 2, 0, 1, 4, 4, 3, 3, 3, 4, 4, 4, 3, 4, 4...
## $ Atr6 <int> 0, 0, 3, 3, 1, 2, 3, 1, 4, 2, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1...
## $ Atr7 <int> 0, 0, 2, 3, 0, 0, 4, 0, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0, 1, 0, 0...
## $ Atr8 <int> 0, 4, 1, 3, 0, 0, 3, 3, 3, 2, 4, 4, 4, 4, 4, 4, 4, 4, 3, 4, 3...
## $ Atr9 <int> 0, 4, 1, 3, 0, 0, 2, 3, 3, 2, 4, 4, 3, 3, 3, 3, 3, 4, 4, 4, 3...
## $ Atr10 <int> 0, 4, 2, 3, 0, 1, 2, 2, 3, 2, 3, 3, 4, 4, 4, 2, 2, 3, 4, 3, 3...
## $ Atr11 <int> 1, 4, 3, 4, 0, 0, 2, 4, 3, 3, 4, 4, 3, 3, 3, 4, 4, 4, 3, 4, 4...
## $ Atr12 <int> 0, 3, 4, 3, 1, 2, 2, 3, 3, 0, 4, 4, 4, 4, 4, 4, 4, 4, 3, 4, 3...
## $ Atr13 <int> 1, 4, 2, 3, 0, 1, 2, 2, 3, 0, 4, 4, 3, 3, 3, 4, 4, 4, 3, 4, 4...
## $ Atr14 <int> 1, 0, 3, 4, 1, 0, 3, 3, 3, 2, 4, 4, 4, 4, 4, 4, 4, 4, 3, 4, 3...
## $ Atr15 <int> 0, 4, 3, 3, 1, 2, 2, 4, 3, 1, 4, 4, 3, 3, 3, 3, 3, 4, 4, 4, 3...
## $ Atr16 <int> 1, 4, 3, 3, 1, 0, 3, 3, 3, 0, 3, 3, 4, 4, 4, 2, 2, 3, 4, 3, 3...
## $ Atr17 <int> 0, 4, 3, 3, 1, 2, 3, 2, 3, 1, 4, 4, 3, 3, 3, 4, 4, 4, 3, 4, 4...
## $ Atr18 <int> 0, 4, 3, 3, 1, 1, 3, 3, 3, 2, 4, 4, 4, 4, 4, 4, 4, 4, 3, 4, 3...
## $ Atr19 <int> 0, 3, 3, 3, 2, 0, 3, 2, 3, 1, 4, 4, 3, 3, 3, 4, 4, 4, 3, 4, 4...
## $ Atr20 <int> 1, 2, 2, 4, 1, 1, 2, 1, 3, 0, 4, 4, 4, 4, 4, 4, 4, 4, 3, 4, 3...
## $ Atr21 <int> 0, 1, 1, 1, 1, 0, 3, 2, 2, 0, 4, 4, 3, 3, 3, 3, 3, 4, 4, 4, 3...
## $ Atr22 <int> 0, 1, 0, 1, 0, 0, 3, 1, 2, 0, 3, 3, 4, 4, 4, 2, 2, 3, 4, 3, 3...
## $ Atr23 <int> 0, 0, 1, 1, 0, 0, 3, 1, 2, 0, 4, 4, 3, 3, 3, 4, 4, 4, 3, 4, 4...
## $ Atr24 <int> 0, 2, 2, 1, 0, 0, 3, 2, 3, 1, 4, 4, 4, 4, 4, 4, 4, 4, 3, 4, 3...
## $ Atr25 <int> 0, 2, 2, 2, 0, 2, 2, 3, 2, 1, 4, 4, 3, 3, 3, 4, 4, 4, 3, 4, 4...
## $ Atr26 <int> 0, 1, 2, 1, 2, 2, 3, 3, 3, 1, 4, 4, 4, 4, 4, 4, 4, 4, 3, 4, 3...
## $ Atr27 <int> 0, 2, 2, 1, 1, 0, 3, 2, 2, 1, 4, 4, 3, 3, 3, 3, 3, 4, 4, 4, 3...
## $ Atr28 <int> 0, 0, 2, 1, 2, 0, 2, 2, 3, 1, 3, 3, 4, 4, 4, 2, 2, 3, 4, 3, 3...
## $ Atr29 <int> 0, 1, 3, 1, 1, 0, 2, 2, 2, 1, 4, 4, 3, 3, 3, 4, 4, 4, 3, 4, 4...
## $ Atr30 <int> 1, 1, 2, 3, 1, 0, 2, 3, 3, 1, 4, 4, 4, 4, 4, 4, 4, 4, 3, 4, 3...
## $ Atr31 <int> 1, 0, 3, 2, 1, 4, 1, 1, 1, 1, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4...
## $ Atr32 <int> 2, 4, 3, 3, 1, 1, 2, 1, 1, 1, 4, 4, 4, 4, 4, 4, 4, 3, 4, 3, 4...
## $ Atr33 <int> 1, 2, 1, 2, 1, 1, 2, 0, 1, 0, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4...
## $ Atr34 <int> 2, 3, 1, 2, 1, 1, 1, 2, 1, 1, 4, 4, 4, 4, 4, 4, 4, 3, 4, 3, 4...
## $ Atr35 <int> 0, 0, 1, 1, 0, 1, 1, 2, 1, 0, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4...
## $ Atr36 <int> 1, 2, 1, 1, 0, 1, 2, 1, 1, 0, 4, 4, 4, 4, 4, 4, 4, 3, 4, 3, 4...
## $ Atr37 <int> 2, 3, 2, 3, 0, 1, 3, 4, 1, 1, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4...
## $ Atr38 <int> 1, 4, 1, 3, 0, 2, 2, 4, 2, 1, 4, 4, 4, 4, 4, 4, 4, 3, 4, 3, 4...
## $ Atr39 <int> 3, 2, 3, 4, 2, 0, 2, 4, 2, 2, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4...
## $ Atr40 <int> 3, 4, 3, 4, 1, 2, 3, 4, 2, 2, 4, 4, 4, 4, 4, 4, 4, 3, 4, 3, 4...
## $ Atr41 <int> 2, 2, 3, 2, 0, 2, 3, 4, 2, 1, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4...
## $ Atr42 <int> 1, 2, 3, 2, 2, 1, 3, 4, 2, 2, 4, 4, 4, 4, 4, 4, 4, 3, 4, 3, 4...
```

```

## $ Atr43 <int> 1, 3, 2, 3, 3, 2, 3, 3, 2, 3, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4...
## $ Atr44 <int> 2, 4, 3, 2, 0, 3, 4, 2, 2, 2, 4, 4, 4, 4, 3, 4, 4, 3, 4, 3, 4...
## $ Atr45 <int> 3, 2, 2, 3, 2, 0, 3, 0, 2, 2, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4...
## $ Atr46 <int> 2, 2, 3, 2, 2, 2, 3, 0, 1, 2, 4, 4, 4, 4, 4, 4, 4, 3, 4, 3, 4...
## $ Atr47 <int> 1, 2, 2, 2, 1, 2, 2, 1, 1, 0, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4...
## $ Atr48 <int> 3, 3, 3, 3, 2, 1, 3, 2, 1, 2, 4, 4, 4, 4, 4, 4, 4, 3, 4, 3, 4...
## $ Atr49 <int> 3, 4, 1, 3, 3, 2, 2, 2, 1, 2, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4...
## $ Atr50 <int> 3, 4, 1, 3, 2, 1, 3, 2, 1, 2, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4...
## $ Atr51 <int> 2, 4, 1, 3, 2, 1, 3, 1, 1, 2, 4, 4, 4, 4, 4, 4, 4, 3, 4, 3, 4...
## $ Atr52 <int> 3, 4, 2, 2, 2, 1, 2, 1, 1, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4...
## $ Atr53 <int> 2, 2, 2, 2, 1, 2, 2, 1, 1, 3, 4, 4, 4, 4, 4, 4, 4, 3, 4, 3, 4...
## $ Atr54 <int> 1, 2, 2, 2, 0, 0, 2, 0, 1, 3, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4...
## $ Class <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1...

```

The survey design has some bias. From the initial article, the researchers explained that Turkey is among the 12 countries with increasing divorce rate (as of 2019 when the study was published). Divorced participants answered based on their marriage experience, while the married participants included were the ones who considered themselves as having a “happy marriage”, without any thought of divorce. This means, married people who are having problems and have signs of divorce were excluded.

The dataset does not include any demographic variables and only have the 54 questions’ responses and marital status(output variable).

Attribute Information

1. When one of us apologizes when our discussions go bad, the issue does not extend.
2. I know we can ignore our differences, even if things get hard sometimes.
3. When we need to, we can take our discussions from the beginning and correct it.
4. When I argue with my spouse, it will eventually work for me to contact him.
5. The time I spent with my spouse is special for us.
6. We don’t have time at home as partners.
7. We are like two strangers who share the same environment at home rather than family.
8. I enjoy our holidays with my spouse.
9. I enjoy traveling with my spouse.
10. My spouse and most of our goals are common.
11. I think that some day, my spouse and I will be in harmony with each other.
12. My spouse and I have similar values regarding personal freedom.
13. My spouse and I have similar entertainment.
14. Most of our goals in regards to people (children, friends, etc.) are the same.
15. My dreams of living are similar and harmonious with those of my spouse.
16. I’m compatible with my spouse about what love should be.
17. I share the same views with my spouse about being happy.
18. My spouse and I have similar ideas about how marriage should be.
19. My spouse and I have similar ideas about how roles should be in marriage.
20. My spouse and I have similar values regarding trust.
21. I know exactly what my spouse likes.
22. I know how my spouse wants to be taken care of when she’s sick.
23. I know my spouse’s favorite food.
24. I can tell you what kind of stress my spouse is having in life.
25. I have knowledge of my spouse’s inner world.
26. I know my spouse’s basic concerns.
27. I know what my spouse’s current sources of stress are.
28. I know my spouse’s hopes and wishes.
29. I know my spouse very well.
30. I know my spouse’s friends and their social relationships.

31. I feel aggressive when I argue with my spouse.
32. When discussing with my spouse, I usually use expressions such as X, Y, Z.
33. I can use negative statements about my spouse's personality during our discussions.
34. I can use offensive expressions during our discussions.
35. I can insult our discussions.
36. I can be humiliating when we argue.
37. My argument with my spouse is not calm.
38. I hate my spouse's way of bringing it up.
39. Fights often occur suddenly.
40. We're just starting a fight before I know what's going on.
41. When I talk to my spouse about something, my calm suddenly breaks.
42. When I argue with my spouse, it only snaps in and I don't say a word.
43. I'm mostly willing to calm the environment a little bit.
44. Sometimes I think it's good for me to leave home for a while.
45. I'd rather stay silent than argue with my spouse.
46. Even if I'm right in the argument, I'm willing not to upset the other side.
47. When I argue with my spouse, I remain silent because I am afraid of not being able to control my anger.
48. I feel right in our discussions.
49. I have nothing to do with what I've been accused of.
50. I'm not actually the one who's guilty of what I'm accused of.
51. I'm not the one who's wrong about problems at home.
52. I wouldn't hesitate to tell her about my spouse's inadequacy.
53. I remind my spouse of her inadequacies during our discussion.
54. I'm not afraid to tell her about my spouse's incompetence.

Data wrangling

The decision variable *Class* is a binary variable, having married as 0 and divorced 1. I will adjust this variable as factor. Other than that, the dataset is very clean.

```
divorce <- divorce %>% dplyr::mutate(Class = factor(Class))
```

Datasets for model creation

This dataset has only 170 rows. Being a small dataset, I will split the data in a 80-20 proportion for training and testing. There will not be a validation dataset.

From the training dataset 50% of participants are divorced. From the testing dataset 48% are divorced. This is a very balanced split.

```
set.seed(1, sample.kind="Rounding")

test_index <- createDataPartition(y = divorce$Class, times = 1, p = 0.2, list = FALSE)

divorce_training <- divorce[-test_index, ]

divorce_testing <- divorce[test_index, ]

print("Training set proportions:")
```

```
## [1] "Training set proportions:"
```

```
prop.table(table(divorce_training$Class))
```

```
##
##           0           1
## 0.5037037 0.4962963
```

```
print("Testing set proportions:")
```

```
## [1] "Testing set proportions:"
```

```
prop.table(table(divorce_testing$Class))
```

```
##
##           0           1
## 0.5142857 0.4857143
```

Data Exploration

- Dataset overview

As explained before, there are 54 predictors, which are answers to a survey. The response scale is (0=Never, 1=Seldom, 2=Averagely, 3=Frequently, 4=Always).

From the data summary we can see that more than 20 answers have a mean greater than 2.

```
summary <- summary(divorce)
summary
```

```
##           Atr1           Atr2           Atr3           Atr4
## Min.      :0.000   Min.      :0.000   Min.      :0.000   Min.      :0.000
## 1st Qu.:0.000   1st Qu.:0.000   1st Qu.:0.000   1st Qu.:0.000
## Median :2.000   Median :2.000   Median :2.000   Median :1.000
## Mean     :1.776   Mean     :1.653   Mean     :1.765   Mean     :1.482
## 3rd Qu.:3.000   3rd Qu.:3.000   3rd Qu.:3.000   3rd Qu.:3.000
## Max.     :4.000   Max.     :4.000   Max.     :4.000   Max.     :4.000
##           Atr5           Atr6           Atr7           Atr8
## Min.      :0.000   Min.      :0.0000   Min.      :0.0000   Min.      :0.000
## 1st Qu.:0.000   1st Qu.:0.0000   1st Qu.:0.0000   1st Qu.:0.000
## Median :1.000   Median :0.0000   Median :0.0000   Median :1.000
## Mean     :1.541   Mean     :0.7471   Mean     :0.4941   Mean     :1.453
## 3rd Qu.:3.000   3rd Qu.:1.0000   3rd Qu.:1.0000   3rd Qu.:3.000
## Max.     :4.000   Max.     :4.0000   Max.     :4.0000   Max.     :4.000
##           Atr9           Atr10          Atr11          Atr12
## Min.      :0.000   Min.      :0.000   Min.      :0.000   Min.      :0.000
## 1st Qu.:0.000   1st Qu.:0.000   1st Qu.:0.000   1st Qu.:0.000
## Median :1.000   Median :2.000   Median :1.000   Median :1.500
## Mean     :1.459   Mean     :1.576   Mean     :1.688   Mean     :1.653
## 3rd Qu.:3.000   3rd Qu.:3.000   3rd Qu.:3.000   3rd Qu.:3.000
## Max.     :4.000   Max.     :4.000   Max.     :4.000   Max.     :4.000
##           Atr13          Atr14          Atr15          Atr16
```

##	Min. :0.000	Min. :0.000	Min. :0.000	Min. :0.000	
##	1st Qu.:0.000	1st Qu.:0.000	1st Qu.:0.000	1st Qu.:0.000	
##	Median :2.000	Median :1.000	Median :1.000	Median :1.000	
##	Mean :1.835	Mean :1.571	Mean :1.571	Mean :1.476	
##	3rd Qu.:3.000	3rd Qu.:3.000	3rd Qu.:3.000	3rd Qu.:3.000	
##	Max. :4.000	Max. :4.000	Max. :4.000	Max. :4.000	
##	Atr17	Atr18	Atr19	Atr20	
##	Min. :0.000	Min. :0.000	Min. :0.000	Min. :0.000	
##	1st Qu.:0.000	1st Qu.:0.000	1st Qu.:0.000	1st Qu.:0.000	
##	Median :1.000	Median :1.000	Median :1.000	Median :1.000	
##	Mean :1.653	Mean :1.518	Mean :1.641	Mean :1.459	
##	3rd Qu.:3.000	3rd Qu.:3.000	3rd Qu.:3.000	3rd Qu.:3.000	
##	Max. :4.000	Max. :4.000	Max. :4.000	Max. :4.000	
##	Atr21	Atr22	Atr23	Atr24	
##	Min. :0.000	Min. :0.000	Min. :0.000	Min. :0.000	
##	1st Qu.:0.000	1st Qu.:0.000	1st Qu.:0.000	1st Qu.:0.000	
##	Median :1.000	Median :0.000	Median :0.000	Median :1.000	
##	Mean :1.388	Mean :1.247	Mean :1.412	Mean :1.512	
##	3rd Qu.:3.000	3rd Qu.:3.000	3rd Qu.:3.000	3rd Qu.:3.000	
##	Max. :4.000	Max. :4.000	Max. :4.000	Max. :4.000	
##	Atr25	Atr26	Atr27	Atr28	Atr29
##	Min. :0.000	Min. :0.000	Min. :0.0	Min. :0.000	Min. :0.000
##	1st Qu.:0.000	1st Qu.:0.000	1st Qu.:0.0	1st Qu.:0.000	1st Qu.:0.000
##	Median :1.000	Median :1.000	Median :1.0	Median :0.500	Median :1.000
##	Mean :1.629	Mean :1.488	Mean :1.4	Mean :1.306	Mean :1.494
##	3rd Qu.:3.000	3rd Qu.:3.000	3rd Qu.:3.0	3rd Qu.:3.000	3rd Qu.:3.000
##	Max. :4.000	Max. :4.000	Max. :4.0	Max. :4.000	Max. :4.000
##	Atr30	Atr31	Atr32	Atr33	Atr34
##	Min. :0.000	Min. :0.000	Min. :0.000	Min. :0.000	Min. :0.0
##	1st Qu.:0.000	1st Qu.:0.000	1st Qu.:0.000	1st Qu.:0.000	1st Qu.:0.0
##	Median :1.000	Median :2.000	Median :2.000	Median :1.000	Median :1.0
##	Mean :1.494	Mean :2.124	Mean :2.059	Mean :1.806	Mean :1.9
##	3rd Qu.:3.000	3rd Qu.:4.000	3rd Qu.:4.000	3rd Qu.:4.000	3rd Qu.:4.0
##	Max. :4.000	Max. :4.000	Max. :4.000	Max. :4.000	Max. :4.0
##	Atr35	Atr36	Atr37	Atr38	
##	Min. :0.000	Min. :0.000	Min. :0.000	Min. :0.000	
##	1st Qu.:0.000	1st Qu.:0.000	1st Qu.:0.000	1st Qu.:0.000	
##	Median :0.500	Median :0.000	Median :2.000	Median :1.000	
##	Mean :1.671	Mean :1.606	Mean :2.088	Mean :1.859	
##	3rd Qu.:4.000	3rd Qu.:4.000	3rd Qu.:4.000	3rd Qu.:4.000	
##	Max. :4.000	Max. :4.000	Max. :4.000	Max. :4.000	
##	Atr39	Atr40	Atr41	Atr42	
##	Min. :0.000	Min. :0.000	Min. :0.000	Min. :0.000	
##	1st Qu.:0.000	1st Qu.:0.000	1st Qu.:0.000	1st Qu.:0.000	
##	Median :2.000	Median :1.500	Median :2.000	Median :2.000	
##	Mean :2.088	Mean :1.871	Mean :1.994	Mean :2.159	
##	3rd Qu.:4.000	3rd Qu.:4.000	3rd Qu.:4.000	3rd Qu.:4.000	
##	Max. :4.000	Max. :4.000	Max. :4.000	Max. :4.000	
##	Atr43	Atr44	Atr45	Atr46	
##	Min. :0.000	Min. :0.000	Min. :0.000	Min. :0.000	
##	1st Qu.:2.000	1st Qu.:0.000	1st Qu.:1.000	1st Qu.:2.000	
##	Median :3.000	Median :2.000	Median :3.000	Median :3.000	
##	Mean :2.706	Mean :1.941	Mean :2.459	Mean :2.553	
##	3rd Qu.:4.000	3rd Qu.:4.000	3rd Qu.:4.000	3rd Qu.:4.000	

```
## Max. :4.000 Max. :4.000 Max. :4.000 Max. :4.000
## Atr47 Atr48 Atr49 Atr50
## Min. :0.000 Min. :0.000 Min. :0.000 Min. :0.000
## 1st Qu.:1.000 1st Qu.:2.000 1st Qu.:1.000 1st Qu.:1.000
## Median :2.000 Median :3.000 Median :3.000 Median :2.000
## Mean :2.271 Mean :2.741 Mean :2.382 Mean :2.429
## 3rd Qu.:4.000 3rd Qu.:4.000 3rd Qu.:4.000 3rd Qu.:4.000
## Max. :4.000 Max. :4.000 Max. :4.000 Max. :4.000
## Atr51 Atr52 Atr53 Atr54 Class
## Min. :0.000 Min. :0.000 Min. :0.000 Min. :0.000 0:86
## 1st Qu.:2.000 1st Qu.:1.000 1st Qu.:1.000 1st Qu.:0.000 1:84
## Median :3.000 Median :3.000 Median :2.000 Median :2.000
## Mean :2.476 Mean :2.518 Mean :2.241 Mean :2.012
## 3rd Qu.:4.000 3rd Qu.:4.000 3rd Qu.:4.000 3rd Qu.:4.000
## Max. :4.000 Max. :4.000 Max. :4.000 Max. :4.000
```

- Missing values

The dataset does not have any missing values that need to be removed or handle.

```
table(is.na(divorce))
```

```
##
## FALSE
## 9350
```

- Histograms of predictor variables

A visual inspection tells us that the attributes that could predict divorce based on the counts and more filled as divorced are: atr1, atr3, atr13, atr25, atr32, atr39.

```
## store predictor variables
attributes <- setdiff(colnames(divorce), "Class")

## initialize an empty list to store plots
histograms = list()

## initialize plot index
i = 1

## create plot and fill plot list
for (attr in attributes) {
  hist = ggplot2::ggplot(divorce_training) +
    ggplot2::geom_histogram(aes_string(x=attr, fill="Class"),
                           position="stack", alpha=0.6) +
    ggplot2::theme_classic() +
    ggplot2::theme(legend.position="bottom")

  histograms[[i]] = hist
  i = i + 1
}
```




```
## display plots. I will plot this graphs at the end after the references for clarity
#histograms
#do.call("grid.arrange", c(histograms, ncol=3))
```

- Evaluating variables correlation

```
main_vars <- divorce_training[, 1:54]

#calculating correlation among variables in divorce
corr_divorce <- cor(main_vars)

corrplot::corrplot(corr_divorce, type="upper", method="shade", tl.cex = 0.6)
```



The correlation chart shows there are many variables highly correlated, as shown with the darkets blue shade (values close to 1). We need to find the relevant variables to run a predictive model, thus, we want to remove highly correlated variables to avoid overfitting.

- Eliminating highly correlated variables

To efficiently reduce the amount of predictors, we will be using Principal Component Analisis (PCA)

The first 9 elements account for 90% of the variability. The 25 first elements account for 98% of the variability. Only the first component explains 77% of the variability.


```
x <- divorce_training[, attributes] #Selecting attributes from dataset

divorce_training_pca <- prcomp(x) #Calculating PC from attributes
summary(divorce_training_pca)
```

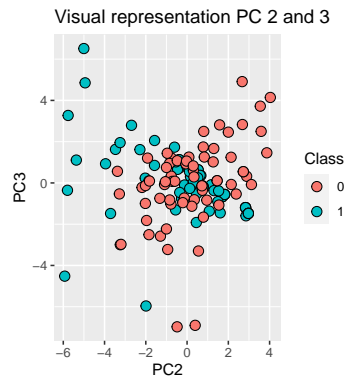
```
## Importance of components:
##          PC1          PC2          PC3          PC4          PC5          PC6          PC7
```

```
## Standard deviation      9.9169 2.10905 1.90120 1.40838 1.32143 1.26750 1.20208
## Proportion of Variance 0.7724 0.03493 0.02839 0.01558 0.01371 0.01262 0.01135
## Cumulative Proportion 0.7724 0.80732 0.83571 0.85129 0.86501 0.87762 0.88897
##          PC8      PC9      PC10      PC11      PC12      PC13      PC14
## Standard deviation      1.12240 1.09812 1.01164 0.97886 0.93897 0.85697 0.83375
## Proportion of Variance 0.00989 0.00947 0.00804 0.00753 0.00692 0.00577 0.00546
## Cumulative Proportion 0.89887 0.90834 0.91637 0.92390 0.93082 0.93659 0.94205
##          PC15      PC16      PC17      PC18      PC19      PC20      PC21
## Standard deviation      0.81812 0.74315 0.72554 0.69914 0.68521 0.61560 0.57906
## Proportion of Variance 0.00526 0.00434 0.00413 0.00384 0.00369 0.00298 0.00263
## Cumulative Proportion 0.94731 0.95165 0.95578 0.95962 0.96331 0.96628 0.96892
##          PC22      PC23      PC24      PC25      PC26      PC27      PC28
## Standard deviation      0.57025 0.54391 0.53732 0.51584 0.50180 0.47519 0.44590
## Proportion of Variance 0.00255 0.00232 0.00227 0.00209 0.00198 0.00177 0.00156
## Cumulative Proportion 0.97147 0.97379 0.97606 0.97815 0.98013 0.98190 0.98346
##          PC29      PC30      PC31      PC32      PC33      PC34      PC35
## Standard deviation      0.43946 0.43497 0.41335 0.38767 0.36893 0.36494 0.35122
## Proportion of Variance 0.00152 0.00149 0.00134 0.00118 0.00107 0.00105 0.00097
## Cumulative Proportion 0.98498 0.98647 0.98781 0.98899 0.99006 0.99110 0.99207
##          PC36      PC37      PC38      PC39      PC40      PC41      PC42
## Standard deviation      0.34239 0.32768 0.3189 0.29698 0.27924 0.26448 0.25400
## Proportion of Variance 0.00092 0.00084 0.0008 0.00069 0.00061 0.00055 0.00051
## Cumulative Proportion 0.99299 0.99384 0.9946 0.99533 0.99594 0.99649 0.99700
##          PC43      PC44      PC45      PC46      PC47      PC48      PC49
## Standard deviation      0.24348 0.2247 0.21123 0.20330 0.19254 0.18120 0.16258
## Proportion of Variance 0.00047 0.0004 0.00035 0.00032 0.00029 0.00026 0.00021
## Cumulative Proportion 0.99746 0.9979 0.99821 0.99853 0.99882 0.99908 0.99929
##          PC50      PC51      PC52      PC53      PC54
## Standard deviation      0.15695 0.14094 0.13272 0.12741 0.10986
## Proportion of Variance 0.00019 0.00016 0.00014 0.00013 0.00009
## Cumulative Proportion 0.99948 0.99964 0.99978 0.99991 1.00000
```

```
data.frame(divorce_training_pca$x[,1:2], Class=divorce_training$Class) %>%
  ggplot2::ggplot(aes(PC1,PC2, fill = Class))+
  ggplot2::geom_point(cex=3, pch=21) +
  ggplot2::coord_fixed(ratio = 1)+
  ggplot2::ggtitle("Visual representation First Two Principal Components")
```



```
data.frame(divorce_training_pca$x[,2:3], Class=divorce_training$Class) %>%
  ggplot2::ggplot(aes(PC2,PC3, fill = Class))+
  ggplot2::geom_point(cex=3, pch=21) +
  ggplot2::coord_fixed(ratio = 1)+
  ggplot2::ggtitle("Visual representation PC 2 and 3")
```



Now we will create a new training and testing dataset with principal components to use in the modeling session.


```
divorce_training_pc = as.data.frame(divorce_training_pca$x)
divorce_training_pc$Class = divorce_training$Class
head(divorce_training_pc)
```

##	PC1	PC2	PC3	PC4	PC5	PC6	PC7
## 1	5.001927	0.4482030	-1.9183911	-0.4050889	2.07916074	-2.4164406	-0.494351
## 3	-2.131388	-3.4694589	1.6256846	2.3380452	-0.78482827	-1.1632855	1.539369
## 4	-4.752328	-3.9502921	0.9283251	0.8547924	0.52184214	-3.3619236	1.761967
## 5	6.056497	-0.6087961	1.7336813	-0.6574911	0.78846621	-0.2286895	-0.505369
## 6	6.095386	-0.5550061	-1.3081149	2.7564768	0.77774284	1.6858438	1.009736
## 7	-4.629296	-2.6933806	2.7921849	0.4245280	0.03805505	-0.8893236	1.111097
##	PC8	PC9	PC10	PC11	PC12	PC13	PC14
## 1	-0.3828507	0.3213686	0.9630951	3.1690760	-1.7297686	0.7022726	-0.09165216
## 3	1.3076764	2.0083988	0.8576773	-0.8056156	0.4054828	-0.1688673	-0.13026595
## 4	0.1093031	-0.4881266	0.9061661	0.2117650	0.1587790	0.1896811	1.14357272
## 5	-0.6234749	-0.3774705	0.4807929	0.9321649	0.5315559	-1.7625150	-0.36416526
## 6	1.1923196	1.6190475	1.2334267	-2.2687860	-0.1356265	0.9835357	1.26326869
## 7	1.9948610	-0.2144704	1.2097839	-0.4789832	-1.3445787	-1.4860855	-1.38843781
##	PC15	PC16	PC17	PC18	PC19	PC20	PC21
## 1	0.8780461	-0.1607495	-1.2302091	0.7919400	0.36003064	-1.1979412	-1.2224653
## 3	1.7475863	-0.1860201	1.5511447	0.4575502	-0.09111096	-0.9224207	-0.4927741
## 4	-0.9600942	1.3454632	0.8642120	2.0205829	0.29516376	-1.0793850	0.3175153
## 5	0.8694334	-0.8452825	1.1500063	0.9531644	-0.43131328	-0.0697104	0.2276545
## 6	-0.5504636	-1.5397232	0.9459050	-1.4566373	0.11364548	-1.3958371	-0.6464744
## 7	0.2923353	0.3649304	-0.9555119	0.1066875	-1.01624678	1.1742153	0.2529383
##	PC22	PC23	PC24	PC25	PC26	PC27	
## 1	-0.127571579	-0.07161427	0.25600716	0.368479171	0.77664993	0.8931457	
## 3	0.003912372	-0.47243435	-0.18952314	-1.001357301	0.20558906	-0.4275691	
## 4	-0.052130053	0.65243890	-0.53103061	0.009104412	0.16427742	0.8605826	
## 5	-0.398490059	-0.36803277	-1.15774333	-0.012233225	1.66455407	-1.5377373	

```
## 6 0.252193479 0.67587544 0.09738649 1.040229824 0.21877661 -1.2836662
## 7 0.299361558 1.32286848 1.80090766 -0.705892002 0.01290306 0.7781087
##      PC28      PC29      PC30      PC31      PC32      PC33
## 1 0.2831462 0.18575733 1.00722010 0.27743906 0.561596537 -0.13752957
## 3 -0.5251967 0.80721182 -0.03782090 -1.73513270 0.687582514 -0.17270602
## 4 1.4692071 -1.23530007 -0.26844745 0.18048194 0.003543012 0.09525928
## 5 0.2890830 0.08404252 1.36860394 0.12308114 -0.640749245 0.11019362
## 6 0.5390956 0.68130435 0.07000135 -0.07103052 -0.303898503 0.37072636
## 7 -1.1739270 -0.15805067 0.81946269 0.06812937 0.762958374 0.84774121
##      PC34      PC35      PC36      PC37      PC38      PC39
## 1 0.1583347 -0.005603301 -0.6540078 -0.3213327 -0.04936921 -0.9146279
## 3 0.1022437 0.305162933 -0.9323578 0.5605979 -0.20728209 0.1951162
## 4 0.3866160 0.496901705 0.5589019 0.3444398 0.22239626 -0.1210227
## 5 0.1986656 -0.175479502 -0.4004802 -0.4990597 -0.33526177 0.2112228
## 6 0.7961522 -0.608458418 0.8994890 -0.7179210 0.02497511 -0.2468381
## 7 0.3746570 -0.265292997 0.6837553 -0.3537728 0.36073001 0.3866860
##      PC40      PC41      PC42      PC43      PC44      PC45
## 1 -0.02966737 0.1178952 0.20361309 -0.217933068 0.2823155 -0.313384785
## 3 0.39829938 -0.5904353 -0.03062511 -0.256609267 -0.2817708 0.198534294
## 4 0.49401989 0.6681529 0.19735677 0.032685234 -0.1299638 -0.007100035
## 5 -0.37936225 0.4833759 -0.36830480 -0.030735760 0.1918664 0.701884264
## 6 -0.27982142 -0.1132089 0.09504793 -0.100302811 0.3339586 -0.655476789
## 7 -0.40690683 -0.1807528 -0.35601455 -0.002339044 0.2880148 0.334591304
##      PC46      PC47      PC48      PC49      PC50      PC51
## 1 0.13014615 -0.07598285 0.05177478 -0.029211762 -0.02074489 0.11460665
## 3 0.60036620 0.17002531 -0.12623810 -0.131477388 -0.13925237 0.09832337
## 4 -0.08448601 0.42703425 -0.05156991 -0.258556486 0.25761210 0.08405680
## 5 -0.23484856 0.11673140 -0.08753785 0.094750062 0.04304332 -0.31857092
## 6 -0.15935518 0.50241346 -0.10830897 0.003326409 -0.04236347 0.03420550
## 7 -0.04713392 -0.07623524 0.04462132 -0.338139890 0.02031832 -0.18400468
##      PC52      PC53      PC54 Class
## 1 -0.45447250 0.01791921 0.006505286 1
## 3 -0.16046745 -0.18603486 -0.045096214 1
## 4 -0.11839101 -0.04804378 -0.056844668 1
## 5 0.21332102 -0.14694180 -0.173686280 1
## 6 -0.06626001 0.02315277 0.113020238 1
## 7 -0.30166746 0.11309212 -0.032099722 1
```

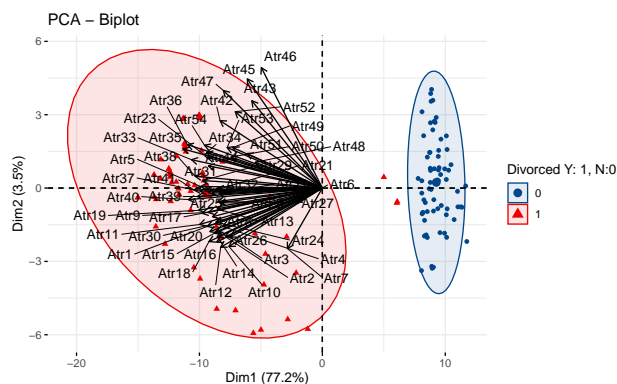
```
divorce_testing_pc = predict(divorce_training_pca, newdata = divorce_testing)
divorce_testing_pc = as.data.frame(divorce_testing_pc)
divorce_testing_pc$class = divorce_testing$class
```

```
factoextra::fviz_eig(divorce_training_pca, addlabels=TRUE,
                      ylim=c(0,80), geom = c("bar", "line"),
                      barfill="lightblue", barcolor="grey", linecolor="blue", ncp=7) +
  labs(title = "Variance Explained By Each Principal Component",
       x = "Principal Components", y = "% of Variance")
```



The following chart also shows how the first component explains 72% of variability.

```
factoextra::fviz_pca_biplot(divorce_training_pca,
                             col.ind = as.factor(divorce_training$Class), col="black",
                             palette = "lancet", geom = c("point"), repel=TRUE,
                             legend.title="Divorced Y: 1, N:0", addEllipses = TRUE)
```



Models

In this section I will explore different methods to predict divorce like logistic regression, kNN, decision trees and random forest and then compare which models provides the best accuracy. I will also use the models in the original data and with the principal components for comparison.

Baseline prediction

Randomly guessing whether a person will get divorced based on the answers to the questionnaire, we get that 51% of them would get divorce. We assume that each person has an equal chance of getting divorced. Randomly guessing gives a divorce estimated greater than the value in the data set which is 43%. The accuracy of randomly guessing is just 63%.

```
set.seed(3, sample.kind = "Rounding")
# guess with equal probability of divorce
guess_divorce <- sample(c(0,1), nrow(divorce_testing), replace = TRUE)
#how many people would be divorced
print("Proportion of people whow would get divorced")
```

```
## [1] "Proportion of people whow would get divorced"

mean(guess_divorce)

## [1] 0.5142857

#Calculating accuracy comparing guessing with testing data set
print("Accuracy")

## [1] "Accuracy"

mean(guess_divorce == divorce_testing$Class)

## [1] 0.6285714

### Accuracy table
summary_accuracy <- tibble(Method = "Baseline Prediction", Accuracy = 0.6285)
kable(summary_accuracy, booktabs = T) %>%
  kableExtra::kable_styling(position = "center", latex_options = "striped")
```

Method	Accuracy
Baseline Prediction	0.6285

Logistic regression

Logistic regression is a regression method where we calculate the probabilities of an output variable to belong to a certain class. Since the logistic regression model provides probabilities between 0-1, to predict the outcome I determine them as 1 or 0 like : if probability > 0.5 the outcome is 1 (divorced, that is our target variable), otherwise it is 0.

- glm model with all data

```
set.seed(3, sample.kind = "Rounding")
logistic_model <- glm(Class ~., family=binomial(link='logit'), data=divorce_training)

#summary(logistic_model)

glm_preds_divorce <- predict(logistic_model, type="response")
print("Divorce rate in logistic regression")

## [1] "Divorce rate in logistic regression"

mean((ifelse(glm_preds_divorce > 0.5, 1, 0)) == divorce_testing$Class)

## [1] 0.5037037
```

- Accuracy in testing set Logistic regression

```

## create prediction probabilities (on test dataset)
glm_test_pred_probs = predict(logistic_model, type="response", newdata=divorce_testing)

## create predictions (on test dataset)
glm_test_preds = as.factor(ifelse(glm_test_pred_probs > 0.5, 1,0))

## evaluate performance (on test dataset)
confusionMatrix(glm_test_preds, divorce_testing$Class)

```

```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0   1
##           0 18   2
##           1   0 15
##
##           Accuracy : 0.9429
##           95% CI : (0.8084, 0.993)
##    No Information Rate : 0.5143
##    P-Value [Acc > NIR] : 4.406e-08
##
##           Kappa : 0.8852
##
##    McNemar's Test P-Value : 0.4795
##
##           Sensitivity : 1.0000
##           Specificity : 0.8824
##           Pos Pred Value : 0.9000
##           Neg Pred Value : 1.0000
##           Prevalence : 0.5143
##           Detection Rate : 0.5143
##    Detection Prevalence : 0.5714
##           Balanced Accuracy : 0.9412
##
##           'Positive' Class : 0
##

```

```

### Accuracy table
# Update table based on testing data
summary_accuracy <- bind_rows(summary_accuracy,
                               tibble(Method = "Logistic Regression",
                                         Accuracy = 0.9429))

# Show table
knitr::kable(summary_accuracy)%>%
  kableExtra::kable_styling(position = "center", latex_options = "striped")

```

Method	Accuracy
Baseline Prediction	0.6285
Logistic Regression	0.9429

- glm model with Principal Components data

```
set.seed(3, sample.kind = "Rounding")
logistic_model_pc <- glm(Class ~.,family=binomial(link='logit'),data=divorce_training_pc)
```

- Accuracy in testing set Principal Components Logistic regression

```
## create prediction probabilities (on test dataset)
glm_test_pred_probs_pc = predict(logistic_model_pc, type="response",
                                newdata=divorce_testing_pc)

## create predictions (on test dataset)
glm_test_preds_pc = as.factor(ifelse(glm_test_pred_probs_pc > 0.5, 1,0))

## evaluate performance (on test dataset)
confusionMatrix(glm_test_preds_pc, divorce_testing_pc$Class)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0   1
##           0 68   0
##           1   0 67
##
##           Accuracy : 1
##           95% CI : (0.973, 1)
##       No Information Rate : 0.5037
##       P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 1
##
##  Mcnemar's Test P-Value : NA
##
##           Sensitivity : 1.0000
##           Specificity : 1.0000
##           Pos Pred Value : 1.0000
##           Neg Pred Value : 1.0000
##           Prevalence : 0.5037
##           Detection Rate : 0.5037
##       Detection Prevalence : 0.5037
##           Balanced Accuracy : 1.0000
##
##           'Positive' Class : 0
##
```

```
### Accuracy table
# Update the accuracy table base on testing data
summary_accuracy <- bind_rows(summary_accuracy,
                              tibble(Method = "Logistic Regression PC",
                                       Accuracy = 1))
# Show the RMSE improvement
knitr::kable(summary_accuracy)%>%
  kableExtra::kable_styling(position = "center", latex_options = "striped")
```


Method	Accuracy
Baseline Prediction	0.6285
Logistic Regression	0.9429
Logistic Regression PC	1.0000

kNN Model

- Model with all the data

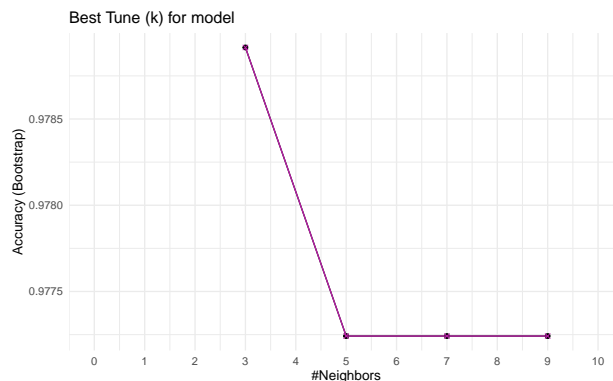
The kNN model (k-nearest neighbors) is a classification algorithm that consists of finding areas called neighbors. It will group the data points in areas measuring the distance between the points. The final value is the value more common among the neighbor.

First, we finding best tune. For this model the best parameter is 3. It means that with 3 neighbors the model can classify the information accurately. Additional neighbors do not improve the accuracy, but actually decreases it.

```
set.seed(3, sample.kind = "Rounding") # simulate R 3.5
train_knn_divorce <- train(Class ~ .,
  method = "knn",
  data = divorce_training,
  tuneGrid = data.frame(k = seq(3, 30, 2)))
train_knn_divorce$bestTune
```

```
##    k
## 1 3
```

```
ggplot2::ggplot(train_knn_divorce) +
  ggplot2::geom_line(colour="#ba1ea8") +
  ggplot2::geom_point(colour="#ba1ea8", shape=4)+
  ggplot2::scale_x_continuous(limits = c(0,10), breaks=seq(0,12,1)) +
  ggplot2::ggtitle("Best Tune (k) for model")+
  theme_minimal()
```



```
knn_preds_divorce <- predict(train_knn_divorce, divorce_testing)
mean(knn_preds_divorce == divorce_testing$Class)
```

```
## [1] 0.9714286
```

```

### Accuracy table
# Update table
summary_accuracy <- bind_rows(summary_accuracy,
                               tibble(Method = "kNN original data",
                                       Accuracy = 0.9714))

# Show table
knitr::kable(summary_accuracy)%>%
  kableExtra::kable_styling(position = "center", latex_options = "striped")

```

Method	Accuracy
Baseline Prediction	0.6285
Logistic Regression	0.9429
Logistic Regression PC	1.0000
kNN original data	0.9714

- kNN Crossvalidation

Now I want to try crossvalidation to see if we can get a better model. Crossvalidation here is a nice tool since the dataset is small.

```

set.seed(8, sample.kind = "Rounding") # simulate R 3.5
train_knn_cv_divorce <- train(Class ~ .,
                              method = "knn",
                              data = divorce_training,
                              tuneGrid = data.frame(k = seq(3, 30, 2)),
                              trControl = trainControl(method = "cv", number = 10, p = 0.9))
train_knn_cv_divorce$bestTune

```

```

##      k
## 14 29

```

```

knn_cv_preds_divorce <- predict(train_knn_cv_divorce, divorce_testing)
mean(knn_cv_preds_divorce == divorce_testing$Class)

```

```
## [1] 0.9714286
```

```

### Accuracy table
# Update the table
summary_accuracy <- bind_rows(summary_accuracy,
                               tibble(Method = "kNN original data Cross Validation",
                                       Accuracy = 0.9714))

# Show table
knitr::kable(summary_accuracy)%>%
  kableExtra::kable_styling(position = "center", latex_options = "striped")

```

Method	Accuracy
Baseline Prediction	0.6285
Logistic Regression	0.9429
Logistic Regression PC	1.0000
kNN original data	0.9714
kNN original data Cross Validation	0.9714

Crossvalidation in this case does not improve the accuracy.

- kNN with only PCA

Now I want to test if I get better results using the data with principal components. For this model, k is the same as the one obtained with the full training data.

```
set.seed(3, sample.kind = "Rounding") # simulate R 3.5
train_knn_divorce_pc <- train(Class ~ .,
  method = "knn",
  data = divorce_training_pc,
  tuneGrid = data.frame(k = seq(3, 30, 2)))
train_knn_divorce_pc$bestTune
```

```
## k
## 1 3
```

```
knn_cv_preds_divorce_pc <- predict(train_knn_divorce_pc, divorce_testing_pc)
mean(knn_cv_preds_divorce_pc == divorce_testing_pc$Class)
```

```
## [1] 0.9777778
```

```
### Accuracy table
# Update the error table
summary_accuracy <- bind_rows(summary_accuracy,
  tibble(Method = "kNN PC",
    Accuracy = 0.9778))
# Show the RMSE improvement
knitr::kable(summary_accuracy)%>%
  kableExtra::kable_styling(position = "center", latex_options = "striped")
```

Method	Accuracy
Baseline Prediction	0.6285
Logistic Regression	0.9429
Logistic Regression PC	1.0000
kNN original data	0.9714
kNN original data Cross Validation	0.9714
kNN PC	0.9778

Classification Tree

- All the data

A classification tree starts by moving into branches, which are observations about the the output variables to predict it. This is an instinctive model that starts splitting the data and evaluates if it contributes to the decision variable. The process is repeated in a recursive manner evaluating all the variables until the new branches does not add value to the outcome variable.

```

train_rpart_divorce = rpart::rpart(Class ~ .,
  data = divorce_training,
  method = 'class',
  control = rpart.control(minsplit=2),
  model = TRUE)

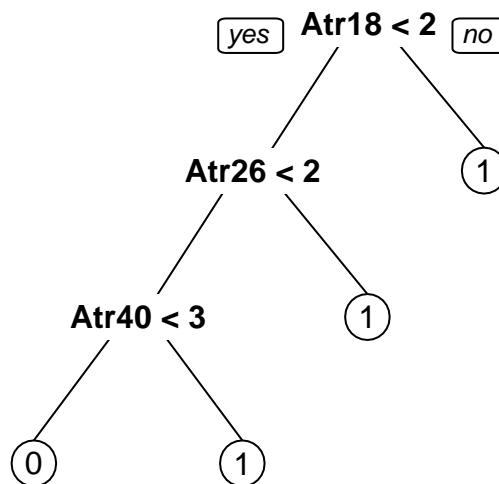
## create prediction probabilities (on train dataset)
rpart_train_pred_probs = predict(train_rpart_divorce)

## create predictions (on train dataset)
rpart_train_preds = as.factor(ifelse(rpart_train_pred_probs[, 2] > 0.5, 1, 0))

rpart.plot::prp(train_rpart_divorce, main="Decision Tree to Predict Divorce")

```

Decision Tree to Predict Divorce



```

ct_preds_divorce <- predict(train_rpart_divorce, divorce_testing)
mean(ct_preds_divorce == divorce_testing$Class)

```

```
## [1] 0.5
```

```

### Accuracy table
# Update the error table
summary_accuracy <- bind_rows(summary_accuracy,
  tibble(Method = "Classification tree",
    Accuracy = 0.5))

```

```
# Show the Accuracy table
knitr::kable(summary_accuracy)%>%
  kableExtra::kable_styling(position = "center", latex_options = "striped")
```

Method	Accuracy
Baseline Prediction	0.6285
Logistic Regression	0.9429
Logistic Regression PC	1.0000
kNN original data	0.9714
kNN original data Cross Validation	0.9714
kNN PC	0.9778
Classification tree	0.5000

```
varImp(train_rpart_divorce)
```

```
##           Overall
## Atr11 59.9407407
## Atr17 61.6425414
## Atr18 61.7498174
## Atr19 61.6425414
## Atr20 61.6872196
## Atr26  3.7754644
## Atr28  1.8607646
## Atr3  0.9710145
## Atr34  0.3043478
## Atr36  2.4140108
## Atr39  0.9710145
## Atr40  3.8317791
## Atr1  0.0000000
## Atr2  0.0000000
## Atr4  0.0000000
## Atr5  0.0000000
## Atr6  0.0000000
## Atr7  0.0000000
## Atr8  0.0000000
## Atr9  0.0000000
## Atr10 0.0000000
## Atr12 0.0000000
## Atr13 0.0000000
## Atr14 0.0000000
## Atr15 0.0000000
## Atr16 0.0000000
## Atr21 0.0000000
## Atr22 0.0000000
## Atr23 0.0000000
## Atr24 0.0000000
## Atr25 0.0000000
## Atr27 0.0000000
## Atr29 0.0000000
## Atr30 0.0000000
## Atr31 0.0000000
## Atr32 0.0000000
```

```
## Atr33 0.0000000
## Atr35 0.0000000
## Atr37 0.0000000
## Atr38 0.0000000
## Atr41 0.0000000
## Atr42 0.0000000
## Atr43 0.0000000
## Atr44 0.0000000
## Atr45 0.0000000
## Atr46 0.0000000
## Atr47 0.0000000
## Atr48 0.0000000
## Atr49 0.0000000
## Atr50 0.0000000
## Atr51 0.0000000
## Atr52 0.0000000
## Atr53 0.0000000
## Atr54 0.0000000
```

From the classification tree we found that the main attributes to predict divorce are Attr 20,19,18,17,11,26,28,3,34,36,39 and 40. All these questions are related to knowing your spouse and her/his needs. When people have low scores on those questions that highly predicts divorce.

- Classification tree principal components


```
train_rpart_divorce_pc = rpart::rpart(Class ~ .,
  data = divorce_training_pc,
  method = 'class',
  control = rpart.control(minsplit=2),
  model = TRUE)

## create prediction probabilities (on train dataset)
rpart_train_pred_probs_pc = predict(train_rpart_divorce_pc)

## create predictions (on train dataset)
rpart_train_preds_pc = as.factor(ifelse(rpart_train_pred_probs_pc[, 2] > 0.5, 1, 0))

rpart.plot::prp(train_rpart_divorce_pc, main="Decision Tree to Predict Divorce PC")
```

Decision Tree to Predict Divorce PC



Since the first principal component explains 73% of the variability, it is not surprised that PC1 is picked on the decision tree. However the accuracy is not improved.

Accuracy

```
ct_preds_divorce_pc <- predict(train_rpart_divorce_pc, divorce_testing_pc)
mean(ct_preds_divorce_pc == divorce_testing_pc$Class)
```

```
## [1] 0.5
```

```
### Accuracy table
# Update the error table
summary_accuracy <- bind_rows(summary_accuracy,
                               tibble(Method = "Classification tree PC",
                                         Accuracy = 0.5))
# Show the Accuracy table
knitr::kable(summary_accuracy)%>%
  kableExtra::kable_styling(position = "center", latex_options = "striped")
```

Method	Accuracy
Baseline Prediction	0.6285
Logistic Regression	0.9429
Logistic Regression PC	1.0000
kNN original data	0.9714
kNN original data Cross Validation	0.9714
kNN PC	0.9778
Classification tree	0.5000
Classification tree PC	0.5000

Random Forest

- All data

```
set.seed(14, sample.kind = "Rounding")    # simulate R 3.5
train_rf_divorce <- train(Class ~ .,
  data = divorce_training,
  method = "rf",
  ntree = 100,
  tuneGrid = data.frame(mtry = seq(1:7)))
train_rf_divorce$bestTune
```

```
##      mtry
## 1      1
```

```
rf_preds_divorce <- predict(train_rf_divorce, divorce_testing)
mean(rf_preds_divorce == divorce_testing$Class)
```

```
## [1] 0.9714286
```

```
varImp(train_rf_divorce)
```

```
## rf variable importance
##
##      only 20 most important variables shown (out of 54)
##
##      Overall
## Atr23 100.00
## Atr40  99.80
## Atr36  80.88
## Atr9   74.12
## Atr11  62.61
## Atr17  60.54
## Atr38  57.16
## Atr44  56.42
## Atr41  54.81
## Atr21  54.48
## Atr50  50.18
## Atr27  48.95
## Atr39  48.30
```



```
## Atr26 43.84
## Atr20 43.04
## Atr54 38.87
## Atr33 38.44
## Atr3 36.36
## Atr10 32.97
## Atr31 31.81
```

Random forest shows the top 3 variables predicting divorce are Atr 23,40,36.


23. I know my spouse's favorite food.
24. I can be humiliating when we argue.
25. We're just starting a fight before I know what's going on.

```
train_rf_divorce$finalModel # inspect final model
```

```
##
## Call:
## randomForest(x = x, y = y, ntree = 100, mtry = param$mtry)
##           Type of random forest: classification
##           Number of trees: 100
## No. of variables tried at each split: 1
##
##           OOB estimate of  error rate: 2.22%
## Confusion matrix:
##      0  1 class.error
## 0 68  0  0.00000000
## 1  3 64  0.04477612
```

```
# make plot of decision tree
plot(train_rf_divorce$finalModel, margin = 0.1)
```

train_rf_divorce\$finalModel



The error graphs shows that even though the algorithm ran 100 trees, the error stays constant from tree number 15.

```
train_rf <- randomForest(Class ~ ., data=divorce_training)

confusionMatrix(predict(train_rf, divorce_testing),
                  divorce_testing$Class)$overall["Accuracy"]
```

```
## Accuracy
## 0.9714286
```

```
### Accuracy table
# Update the error table
summary_accuracy <- bind_rows(summary_accuracy,
                              tibble(Method = "Random Forest",
                                       Accuracy = 0.9714))

# Show table
knitr::kable(summary_accuracy)%>%
  kableExtra::kable_styling(position = "center", latex_options = "striped")
```

Method	Accuracy
Baseline Prediction	0.6285
Logistic Regression	0.9429
Logistic Regression PC	1.0000
kNN original data	0.9714
kNN original data Cross Validation	0.9714
kNN PC	0.9778
Classification tree	0.5000
Classification tree PC	0.5000
Random Forest	0.9714

- Random Forest with Principal Components


```
set.seed(14, sample.kind = "Rounding") # simulate R 3.5
train_rf_divorce_pc <- train(Class ~ .,
  data = divorce_training_pc,
  method = "rf",
  ntree = 100,
  tuneGrid = data.frame(mtry = seq(1:7)))
train_rf_divorce$bestTune
```

```
## mtry
## 1 1
```

```
rf_preds_divorce_pc <- predict(train_rf_divorce_pc, divorce_testing_pc)
mean(rf_preds_divorce_pc == divorce_testing_pc$Class)
```

```
## [1] 1
```

```
varImp(train_rf_divorce_pc)$importance %>%
  as.data.frame() %>%
  rownames_to_column() %>%
  arrange(Overall) %>%
  mutate(rowname = forcats::fct_inorder(rowname)) %>%
  ggplot()+
  geom_col(aes(x = rowname, y = Overall))+
  coord_flip()+
  theme_bw()
```




```
train_rf_divorce_pc$finalModel # inspect final model
```

```
##
## Call:
## randomForest(x = x, y = y, ntree = 100, mtry = param$mtry)
##           Type of random forest: classification
##           Number of trees: 100
## No. of variables tried at each split: 7
##
##           OOB estimate of  error rate: 5.19%
## Confusion matrix:
##      0  1 class.error
## 0 68  0  0.0000000
## 1  7 60  0.1044776
```

```
# make plot of decision tree
plot(train_rf_divorce_pc$finalModel, margin = 0.1)
```

train_rf_divorce_pc\$finalModel



```
#text(train_rf_divorce$finalModel)
```

The error graphs shows that even though the algorithm ran 100 trees, the error stays constant after tree number 20.

```
train_rf_pc <- randomForest(Class ~ ., data=divorce_training_pc)
```

```
confusionMatrix(predict(train_rf_pc, divorce_testing_pc),
  divorce_testing_pc$Class)$overall["Accuracy"]
```

```
## Accuracy
##      1
```

```
### Accuracy table
```

```
# Update the error table
```

```
summary_accuracy <- bind_rows(summary_accuracy,
  tibble(Method = "Random Forest PC",
    Accuracy = 1))
```

```
# Show the RMSE improvement
```

```
knitr::kable(summary_accuracy)%>%
```

```
  kableExtra::kable_styling(position = "center", latex_options = "striped")
```

Method	Accuracy
Baseline Prediction	0.6285
Logistic Regression	0.9429
Logistic Regression PC	1.0000
kNN original data	0.9714
kNN original data Cross Validation	0.9714
kNN PC	0.9778
Classification tree	0.5000
Classification tree PC	0.5000
Random Forest	0.9714
Random Forest PC	1.0000

Results

After running logistic regression, kNN method, classification tree and Random Forest on the divorce dataset to predict divorce in couples based on their answers to the survey, I found that the best models are Logistic Regression and Random Forest when using Principal Components in both with Accuracy of 1. The worse model was classification tree. It seems to be because once it takes one branch, adding other variables are not very powerful when predicting in the test data. Random Forest and kNN with the original data have the same accuracy.

Even though this is a very subjective topic, it is interesting that from random forest the main variable to predict divorce is not knowing something basic as your spouse's favorite food. The other 2 variables are related to unhealthy behaviors on the relationship.

Main variables that predict divorce based on Random Forest with all data

23. I know my spouse's favorite food.
24. I can be humiliating when we argue.
25. We're just starting a fight before I know what's going on.

```
knitr::kable(summary_accuracy)%>%
  kableExtra::kable_styling(position = "center", latex_options = "striped")
```

Method	Accuracy
Baseline Prediction	0.6285
Logistic Regression	0.9429
Logistic Regression PC	1.0000
kNN original data	0.9714
kNN original data Cross Validation	0.9714
kNN PC	0.9778
Classification tree	0.5000
Classification tree PC	0.5000
Random Forest	0.9714
Random Forest PC	1.0000

Conclusion

This was a great exercise to compare different classification methods. I wanted to explore this type of models since this is usually applied for situations that are more difficult to represent in a model. In this case, evaluating whether a person would get divorce is very subjective and it is a field out of my expertise.

Just starting, the data exploration requires creativity, We can't have regular histograms or counts like what you do with continuous variables. An additional issue, is that we have more than 50 predictors. Trying to understand all of them was not easy. It was helpful to use the *for* loop.

On the modeling side, it was useful to run the models to see how much the accuracy changes if you use the original data or principal components. Regarding Principal Components, since in this dataset the first component explained more than 70% of the variability it was not surprising that this was the best predictor.

References


1. Divorce Analysis in R by Howard, https://www.rpubs.com/howard_song/538809
2. Data Science Book from Rafael Irizarry, <https://rafalab.github.io/dsbook/>
3. Predicting whether a couple is going to get divorced or not using artificial neural networks by Ibrahim M. Nasser: <http://dstore.alazhar.edu.ps/xmlui/bitstream/handle/123456789/545/IJEAIS191007.pdf?sequence=1&isAllowed=y>
4. Initial research by Mustafa Kemal, Kemal Adem, Tahsin Ilhan and Serhat Kilicarslan: Divorce prediction using correlation based feature selection and artificial neural networks <https://dergipark.org.tr/en/download/article-file/748448>

Others

Dataset variables histograms


```
histograms
```

```
## [[1]]
```

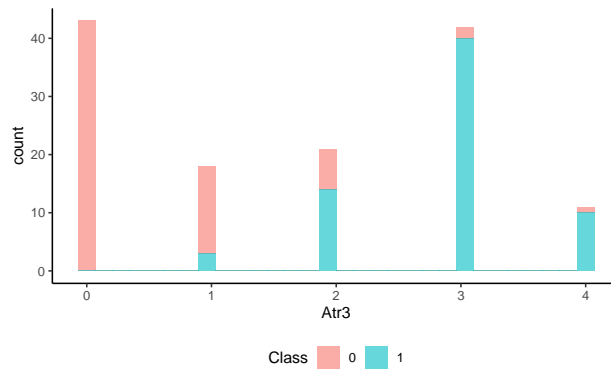


```
##
```


```
## [[2]]
```




[[3]]




[[4]]




[[5]]




[[6]]




[[7]]




[[8]]




```
##
## [[9]]
```




```
##
## [[10]]
```




```
##
## [[11]]
```



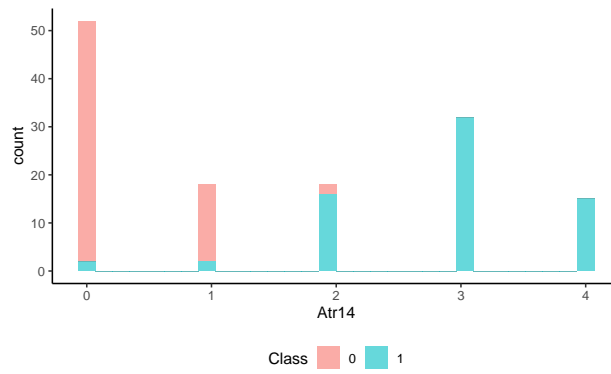
```
##
## [[12]]
```




```
##
## [[13]]
```



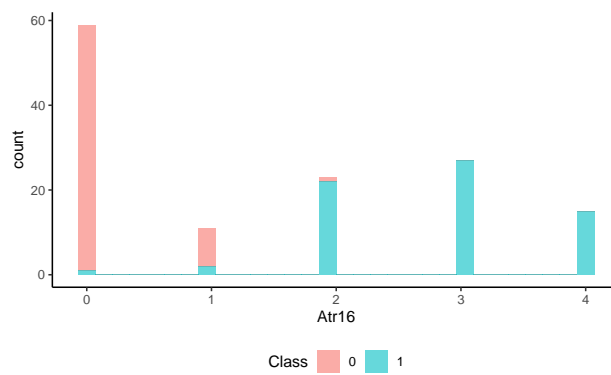
```
##
## [[14]]
```




```
##
## [[15]]
```




```
##
## [[16]]
```



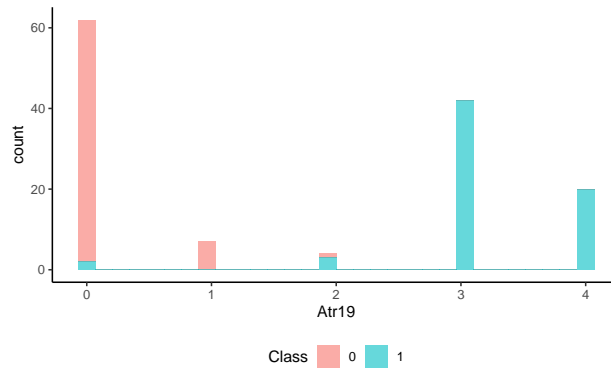
```
##
## [[17]]
```




[[18]]




[[19]]




[[20]]




[[21]]




[[22]]




[[23]]




[[24]]




[[25]]




[[26]]




[[27]]




[[28]]




[[29]]




[[30]]




[[31]]




[[32]]




[[33]]




[[34]]




[[35]]




[[36]]




[[37]]




[[38]]




[[39]]




[[40]]




[[41]]




[[42]]




[[43]]




[[44]]




[[45]]




[[46]]




[[47]]




[[48]]




[[49]]




[[50]]




[[51]]



[[52]]



[[53]]



```
##  
## [[54]]
```

